

Optimizing Compilers

Vineeth Kashyap

Department of Computer Science, UCSB

SIAM Algorithms Seminar, 2014

Warning!!

“Optimizing” compiler is a misnomer

My Background

- Worked on a compiler team for 2 years at National Instruments R&D
- Worked on multiple compilers during my PhD for languages we designed
- Will be working on “compiler-related” research at GrammaTech Research starting January 2015

THE #1 PROGRAMMER EXCUSE
FOR LEGITIMATELY SLACKING OFF:
"MY CODE'S COMPILING."



Credit to xkcd.com

The Beginnings

[. . . .] Programmers wrote assembly

[1952] Grace Hopper wrote the first compiler

[1957] John W. Backus created an optimizing compiler

The Success of Fortran Optimizing Compiler

- Programs took hours instead of weeks to write
- Programs became portable!
- Decisively won the battle against assembly
- Space and military projects used it extensively

FizzBuzz

```
; Init the registers
mov dx,03030h    ; For easier printing, the number is
                  ; kept in Binary Coded Decimal, in
----
;the DX register.
mov ah,0Eh       ; 0Eh is the IBM PC interrupt 10h
                  ;function that does write text on
                  ;the screen in teletype mode.
mov bl,100d      ; BL is the counter (100 numbers).
xor cx,cx        ; CX is a counter that will be used
                  ;for screen printing.
xor bh,bh        ; BH is the counter for counting
                  ;multiples of three.

writeloop:       ; Increment the BCD number in DX.
inc dl           ; Increment the low digit
cmp dl,9Ah       ; If it does not overflow nine,
jnz writeloop1   ;continue with the program,
mov dl,30h       ;otherwise reset it to zero and
inc dh           ;increment the high digit
writeloop1:
inc bh           ; Increment the BH counter.
cmp bh,03h       ; If it reached three, we did
                  ;increment the number three times
                  ;from the last time the number was
                  ;a multiple of three, so the number
                  ;is now a multiple of three now,
                  ;then we need to write "fizz" on the
                  ;screen.
jz writefizz
cmp dl,30h       ; The number isn't a multiple of
jz writebuzz     ;three, so we check if it's a
cmp dl,35h       ;multiple of five. If it is, we
jz writebuzz     ;need to write "buzz". The program
                  ;checks if the last digit is zero or
                  ;five.
mov al,dh        ; If we're here, there's no need to
int 10h          ;write neither "fizz" nor "buzz", so
mov al,dl        ;the program writes the BCD number
int 10h          ;in DX
writespace:
mov al,020h      ;and a white space.
int 10h
dec bl           ; Loop if we didn't process 100
jnz writeloop    ;numbers.

programend:      ; When we did reach 100 numbers,
cli              ;the program flow falls here, where
hlt              ;interrupts are cleared and the
jmp programend   ;program is stopped.

writefizz:       ; There's need to write "fizz":
mov si,offset fizz ; SI points to the "fizz" string,
call write       ;that is written on the screen.
xor bh,bh       ; BH, the counter for computing the
                  ;multiples of three, is cleared.
cmp dl,30h       ; We did write "fizz", but, if the
jz writebuzz     ;number is a multiple of five, we
cmp dl,35h       ;could need to write "buzz" also:
jnz writespace   ;check if the number is multiple of
                  ;five. If not, write a space and
                  ;return to the main loop.
writebuzz:       ; (The above code falls here if
                  ;the last digit is five, otherwise
                  ;it jumps)
mov si,offset buzz ;SI points to the "buzz" string,
call write       ;that is written on the screen.
jmp writespace   ; Write a space to return to the main
                  ;loop.

write:           ; Write subroutine:
mov cl,04h       ; Set CX to the length of the string:
                  ;both strings are 4 bytes long.
writel:
mov al,[si]      ; Load the character to write in AL.
inc si          ; Increment the counter SI.
int 10h          ; Call interrupt 10h, function 0Eh to
                  ;write the character and advance the
                  ;text cursor (teletype mode)
loop writel      ; Decrement CX: if CX is not zero, do
ret              ;loop, otherwise return from
                  ;subroutine.

fizz:            ;The "fizz" string.
db "fizz"

buzz:            ;The "buzz" string.
db "buzz"
```

Credit to Rosetta Code

FizzBuzz

```
; Init the registers
mov dx,03030h    ; For easier printing, the number is
                  ; kept in Binary Coded Decimal, in
;-----
;the DX register.
mov ah,02h       ; 02h is the IBM PC interrupt 10h
                  ;function that does write text on
                  ;the screen in teletype mode.
mov bl,100d      ; BL is the counter (100 numbers).
xor cx,cx        ; CX is a counter that will be used
                  ;for screen printing.
xor bh,bh        ; BH is the counter for counting
                  ;multiples of three.

writeloop:       ; Increment the BCD number in DX.
inc dl           ; Increment the low digit
cmp dl,9ah       ; If it does not overflow nine,
jnz writeloop1   ;continue with the program,
mov dl,30h       ;otherwise reset it to zero and
inc dh           ;increment the high digit
writeloop1:      ; Increment the BH counter.
inc bh          ; If it reached three, we did
cmp bh,03h       ;increment the number three times
                  ;from the last time the number was
                  ;a multiple of three, so the number
                  ;is now a multiple of three now,
                  ;then we need to write "fizz" on the
                  ;screen.
jz writefizz     ; The number isn't a multiple of
                  ;three, so we check if it's a
cmp dl,30h       ;multiple of five. If it is, we
jz writebuzz     ;need to write "buzz". The program
cmp dl,35h       ;checks if the last digit is zero or
jz writebuzz     ;five.
; If we're here, there's no need to
mov al,dh        ;write neither "fizz" nor "buzz", so
int 10h          ;the program writes the BCD number
mov al,dl        ;in DX
int 10h
writespace:      ;and a white space.
mov al,020h
int 10h
dec bl          ; Loop if we didn't process 100
jnz writeloop    ;numbers.

programend:      ; When we did reach 100 numbers,
cli             ;the program flow falls here, where
hlt             ;interrupts are cleared and the
jmp programend   ;program is stopped.

writefizz:       ; There's need to write "fizz":
mov si,offset fizz ; SI points to the "fizz" string,
call write       ;that is written on the screen.
xor bh,bh       ; BH, the counter for computing the
                  ;multiples of three, is cleared.
; We did write "fizz", but, if the
cmp dl,30h       ;number is a multiple of five, we
jz writebuzz     ;could need to write "buzz" also:
cmp dl,35h       ;check if the number is multiple of
jnz writespace   ;five. If not, write a space and
                  ;return to the main loop.
; (The above code falls here if
;the last digit is five, otherwise
;it jumps)
writebuzz:       ;SI points to the "buzz" string,
mov si,offset buzz ;that is written on the screen.
call write       ; Write a space to return to the main
;loop.

write:           ; Write subroutine:
mov cl,04h       ; Set CX to the length of the string:
                  ;both strings are 4 bytes long.
writel:         ; Load the character to write in AL.
mov al,[si]      ; Increment the counter SI.
inc si
int 10h          ; Call interrupt 10h, function 02h to
                  ;write the character and advance the
                  ;text cursor (teletype mode)
loop writel      ; Decrement CX: if CX is not zero, do
ret             ;loop, otherwise return from
                  ;subroutine.

fizz:           ;The "fizz" string.
db "fizz"

buzz:           ;The "buzz" string.
db "buzz"
```

```
program fizzbuzz_if
integer :: i

do i = 1, 100
    if (mod(i,15) == 0) then; print *, 'FizzBuzz'
    else if (mod(i,3) == 0) then; print *, 'Fizz'
    else if (mod(i,5) == 0) then; print *, 'Buzz'
    else;
        print *, i
    end if
end do
end program fizzbuzz_if
```


Credit to Rosetta Code

Compilers 101


Source
Code

Compilers 101

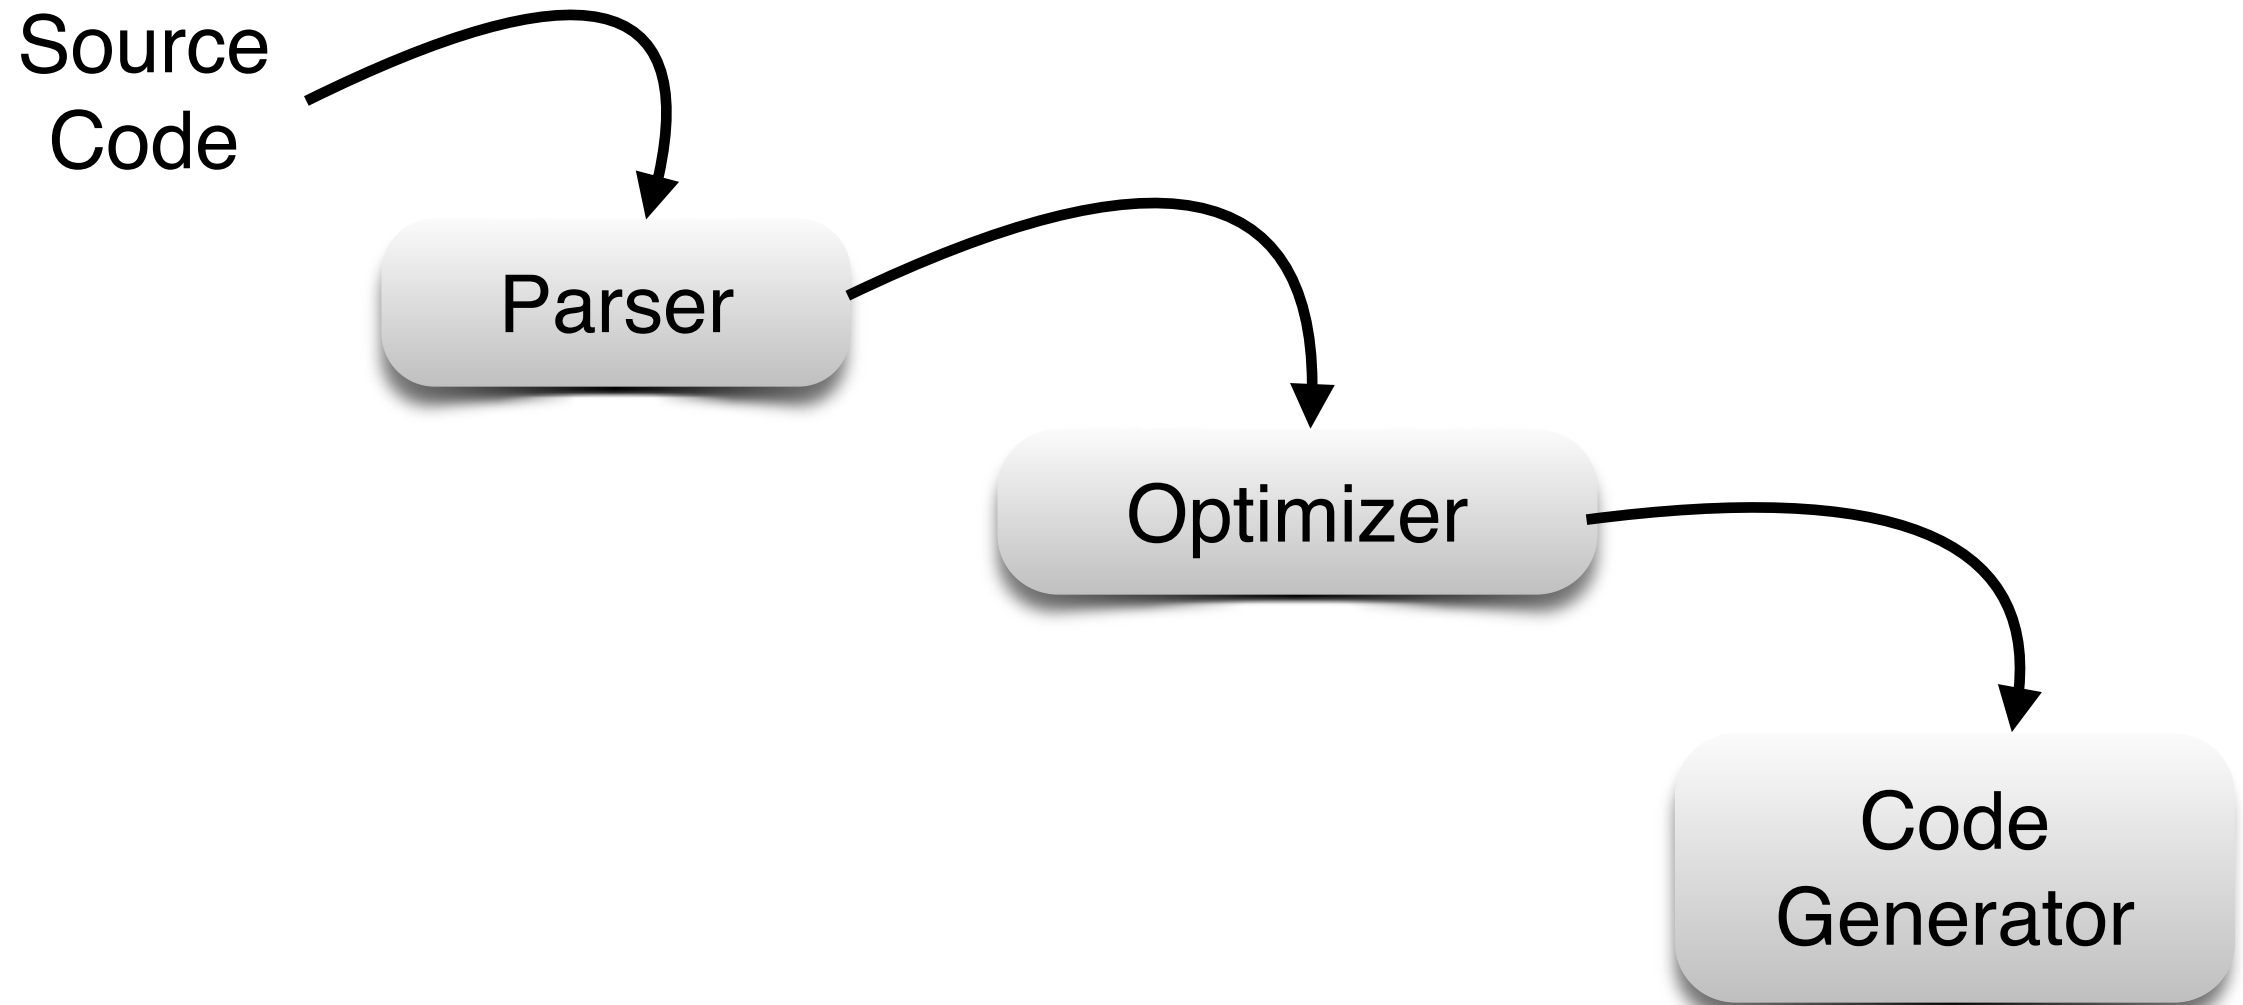
Source
Code



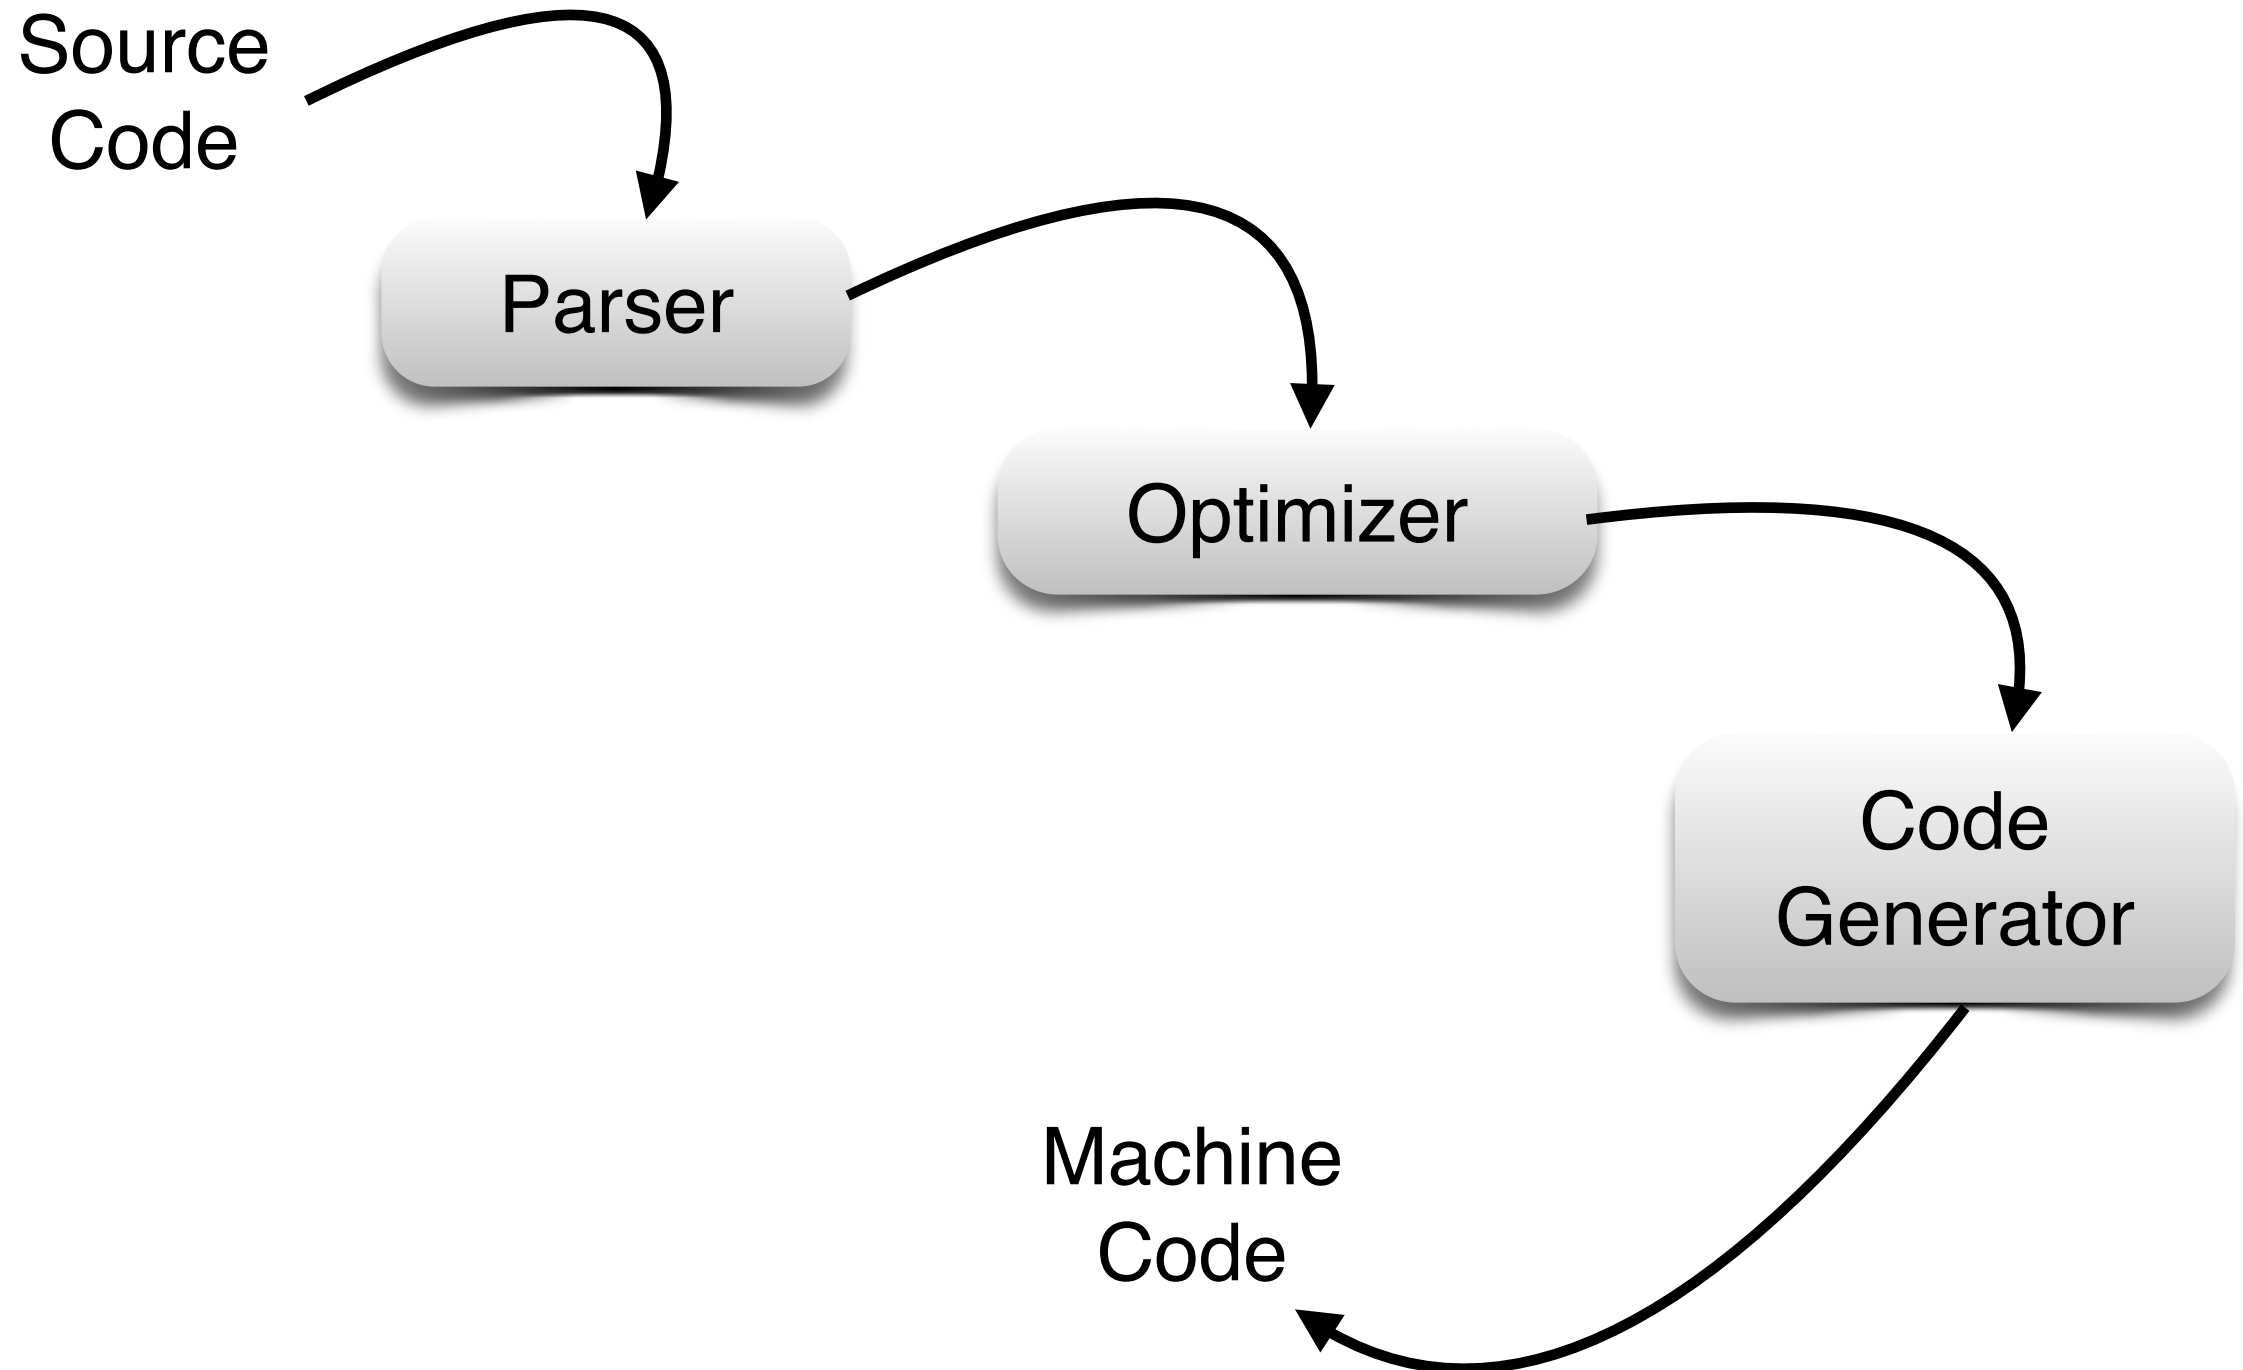
Compilers 101



Compilers 101



Compilers 101




related.classes@cs.ucsb

- Compilers (CS 160)
- Programming Languages (CS 162)
- Program Analysis (CS 260)
- Modern Programming Languages (CS 263)

Rest of the Talk

- A brief bit on parsers
- Data flow analysis
- Some interesting high-level optimizations
- One low level optimization
- A couple of research ideas




Parser

- Converts text (usually) to a structured representation
- Also recognizes a formal language
 - Regular languages (regular expressions)
 - Deterministic context-free grammars

HTML cannot be parsed using regular expressions

Regular expressions can only match regular languages but HTML is a context-free language



Compiler Optimizations

- Finding *equivalent* programs that are *hopefully* better in some way
- It is undecidable what effect an optimization has on performance of a program
- It is undecidable whether an optimization is applicable!
- We opt for safety instead of missed optimization opportunity

Data Flow Analysis

- Figure out how a program manipulates its data
- Ex: “Does a variable always hold the same constant value at a given program point?”
- Ex: “Where are the given definition of a variable used?”
- Important enabler for a lot of optimizations


Conservative Analysis

- Over-approximation: can an optimization be safely performed? (no is acceptable, unless absolutely sure)
- False positives rather than false negatives
- Finds program invariants that hold across all possible executions of the program



Data Flow Analysis Primer

- Operates on elements of a lattice algebraic structure
- Solves simultaneous data flow equations
- Equations involve monotone computations over lattice elements
- Can be solved iteratively, guaranteed to terminate



Constant Propagation



Constant Propagation



Constant Propagation



$$f_{out}(n) = f(T(n), \{f_{out}(k) \mid k \in pred(n)\})$$

Abstract Interpretation


- Can do whatever data flow analysis can do
- Is theoretically elegant, can provide proof of safety
- Mathematically relates the analysis with all executions
- What I use in my dissertation

Loop Invariant Code Motion

```
while (...) {  
    constant := x/y;  
    // other irrelevant code  
}
```

Loop Invariant Code Motion


```
while (...) {  
    constant := x/y;  
    // other irrelevant code  
}
```



```
constant := x/y;  
while (...) {  
    // other irrelevant code  
}
```

Loop Invariant Code Motion

```
while (...) {  
    constant := x/y;  
    // other irrelevant code  
}
```



```
if (...) constant := x/y;  
while (...) {  
    // other irrelevant code  
}
```

Dead Code Elimination

```
var x := 0;  
var y := 0;  
if (x > 0) {  
    // dead code  
}  
print x;
```

Tail Call Optimization

- Turn recursive calls into loops!
- Avoids costs associated with function calls and stack allocations

What are Tail Calls?

```
int sumToN(int n) {  
    if (n == 1)  
        return 1;  
    else  
        return n + sum(n-1);  
} // sumToN(100);
```


What are Tail Calls?

```
int sumToN(int n) {  
    if (n == 1)  
        return 1;  
    else  
        return n + sum(n-1);  
} // sumToN(100);
```

Not a tail call

What are Tail Calls?

```
int sumToN(int n) {  
    if (n == 1)  
        return 1;  
    else  
        return n + sum(n-1);  
} // sumToN(100);
```

Not a tail call

```
int sumToN(int n, int acc) {  
    if (n == 1)  
        return acc;  
    else  
        return sum(n-1, acc+n);  
} // sumToN(100, 1);
```

What are Tail Calls?

```
int sumToN(int n) {  
    if (n == 1)  
        return 1;  
    else  
        return n + sum(n-1);  
} // sumToN(100);
```

Not a tail call

Tail call

```
int sumToN(int n, int acc) {  
    if (n == 1)  
        return acc;  
    else  
        return sum(n-1, acc+n);  
} // sumToN(100, 1);
```

Tail Call Optimized Code

```
int sumToN(int n, int acc) {  
    if (n == 1)  
        return acc;  
    else  
        return sum(n-1, acc+n);  
} // sumToN(100, 1);
```

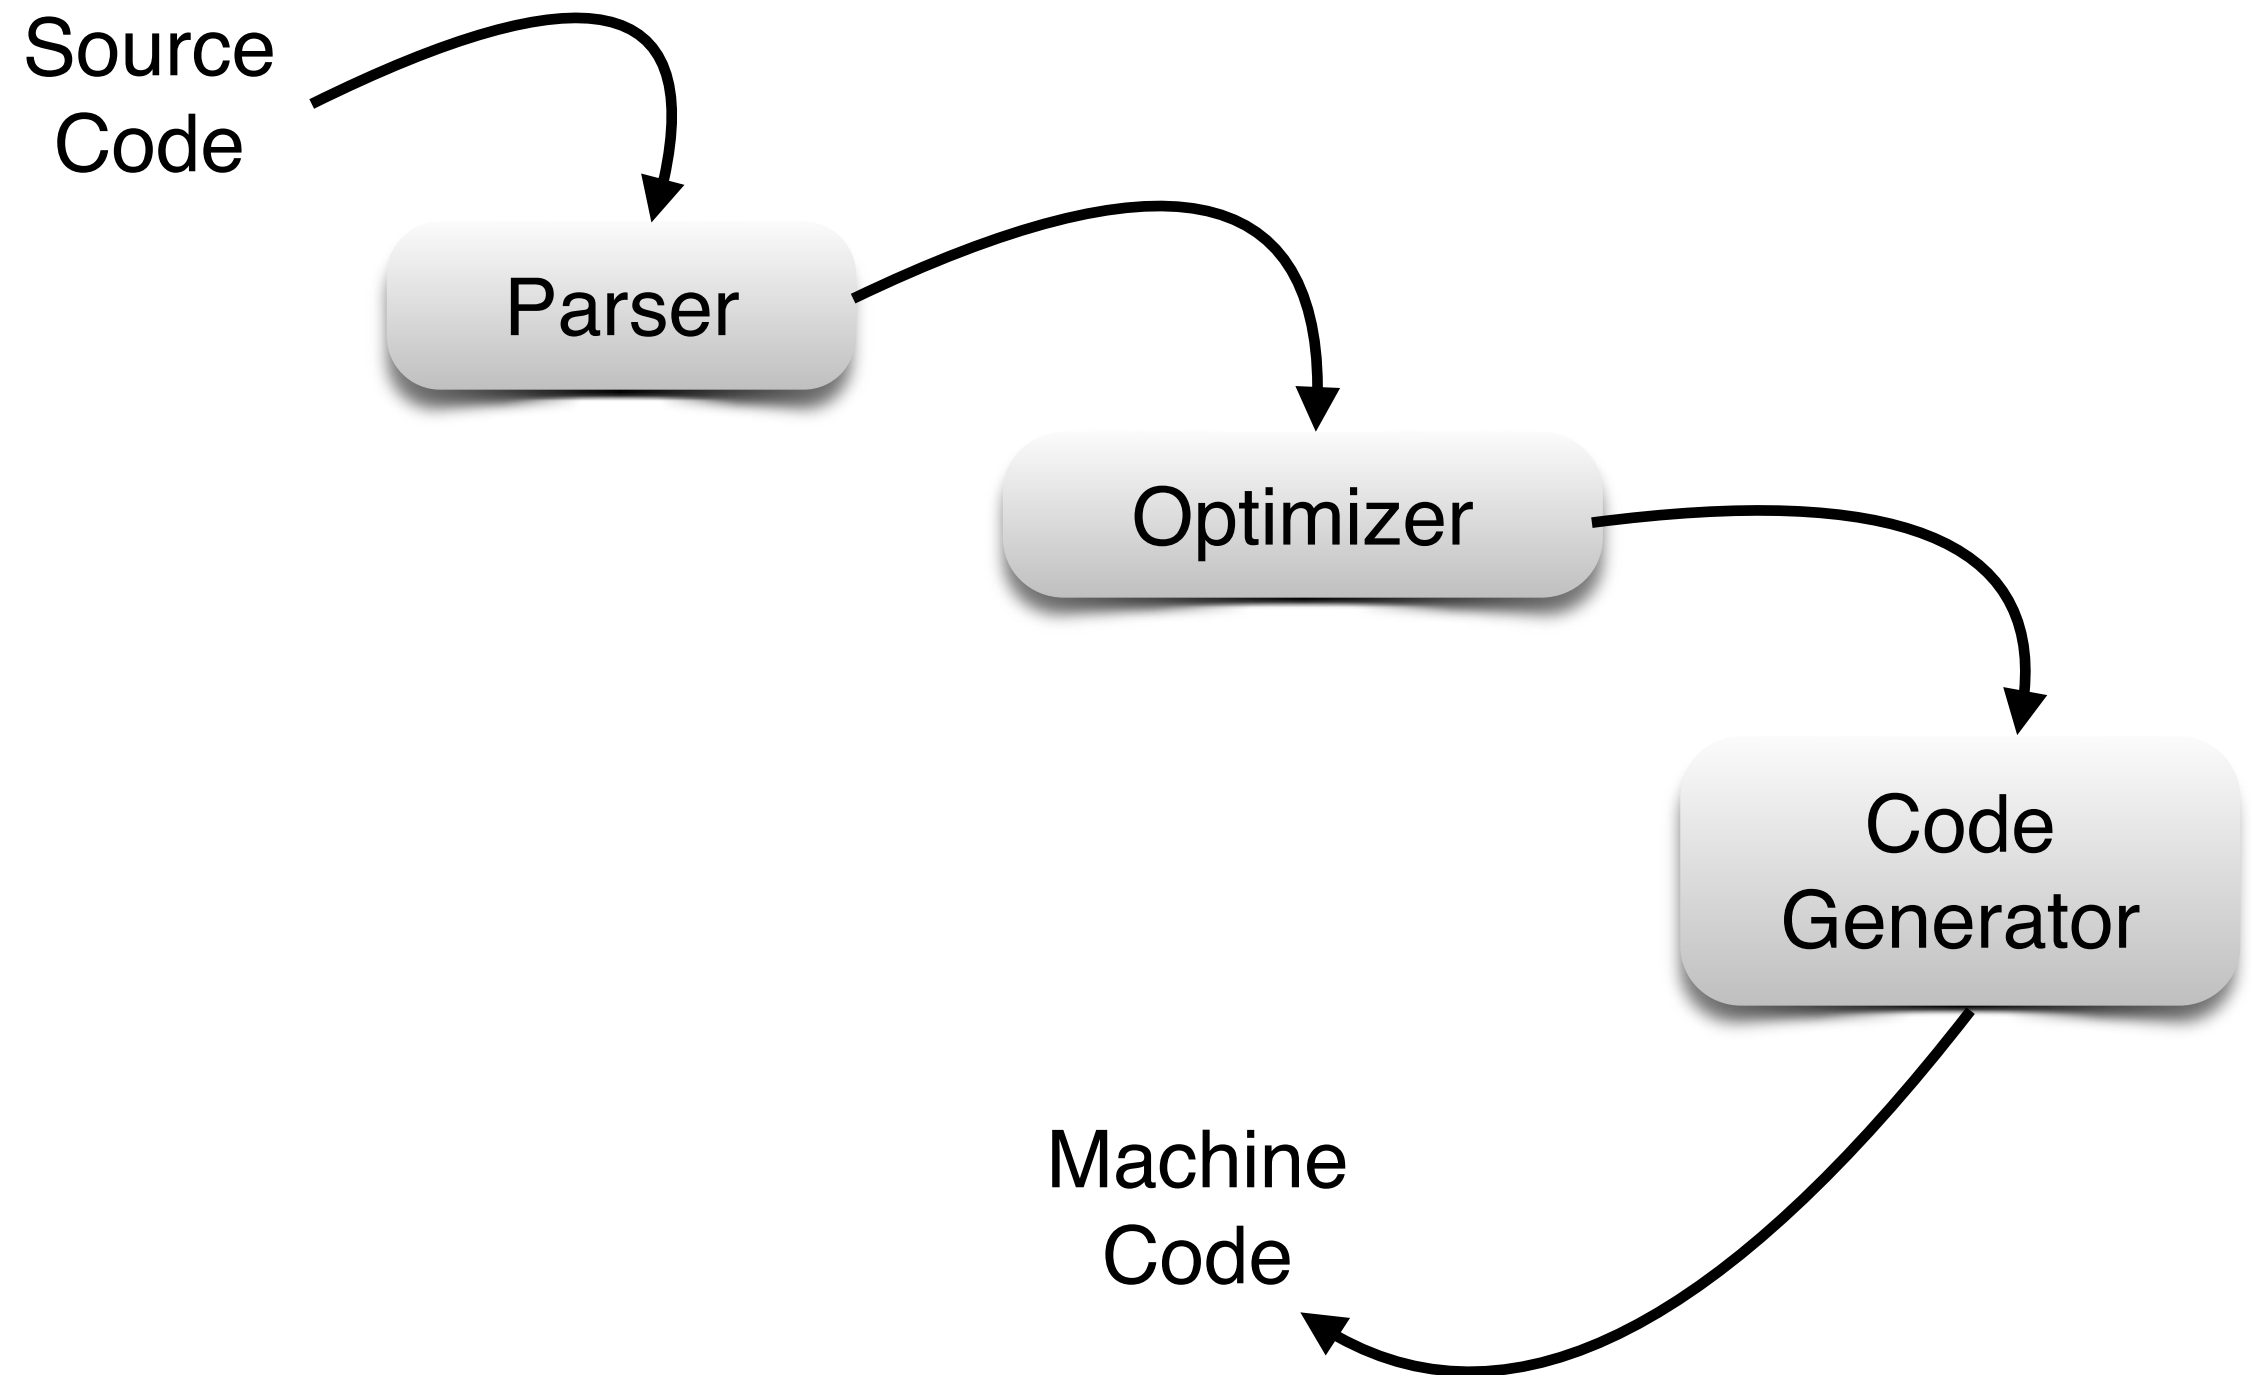
```
int sumToN(int n, int acc) {  
    BEGIN:  
        if (n == 1)  
            return acc;  
        else {  
            acc = acc + n;  
            n = n - 1;  
            goto BEGIN;  
        }  
}
```

Tail Call Optimized Code

```
int sumToN(int n, int acc) {  
    if (n == 1)  
        return acc;  
    else  
        return sum(n-1, acc+n);  
} // sumToN(100, 1);
```

```
int sumToN(int n, int acc) {  
    BEGIN:  
        if (n == 1)  
            return acc;  
        else {  
            acc = acc + n;  
            n = n - 1;  
            goto BEGIN;  
        }  
}
```

Reassign
Parameters



Register Allocation

- Considered the most important optimization
- Minimize traffic between CPU registers and memory
- Usually requires a low level representation

What is Register Allocation

- Determine which values should be in which register
- Registers are scarce resource
- RISC: almost all operations are register based
- CISC: register operations are much faster

Note on Register Assignment

- RISC: usually trivial
- CISC: must take into account special registers

Using Graph Coloring!


- Register allocation can be viewed as graph-coloring [Cocke1971]
- Designed and implemented at IBM [Chaitin1981]

Example Allocation

x	:=	2		
y	:=	4		
w	:=	x	+	y
z	:=	x	+	1
u	:=	x	*	y
x	:=	z	*	2


Symbolic Registers

x	:=	2	
y	:=	4	
w	:=	x + y	
z	:=	x + 1	
u	:=	x * y	
x	:=	z * 2	




s1	:=	2	
s2	:=	4	
s3	:=	s1 + s2	
s4	:=	s1 + 1	
s5	:=	s1 * s2	
s6	:=	s4 * 2	

Interference Graphs



s1	:=	2
s2	:=	4
s3	:=	s1 + s2
s4	:=	s1 + 1
s5	:=	s1 * s2
s6	:=	s4 * 2


Graph Coloring




s1	:=	2
s2	:=	4
s3	:=	s1 + s2
s4	:=	s1 + 1
s5	:=	s1 * s2
s6	:=	s4 * 2

Allocate Registers

```
s1 := 2
s2 := 4
s3 := s1 + s2
s4 := s1 + 1
s5 := s1 * s2
s6 := s4 * 2
```



```
r1 := 2
r2 := 4
r3 := r1 + r2
r3 := r1 + 1
r1 := r1 * r2
r2 := r3 * 2
```



Register Spilling

What if the graph has a chromatic number $k >$ number of registers?

Interesting Research Ideas

- Superoptimization
- Optimizations for reducing energy consumption

Advanced
COMPILER & DESIGN
IMPLEMENTATION

Steven S. Muchnick



Ping me!

Twitter: @vineethkashyap

Email: vineeth@cs.ucsb.edu

Website: <http://www.cs.ucsb.edu/~vineeth/>